



Task 1: DPO conditions

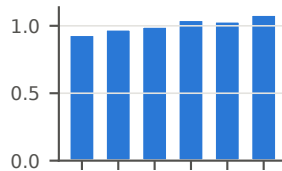
held-out pref. accuracy



KL to reference



reward-model score



response tokens

