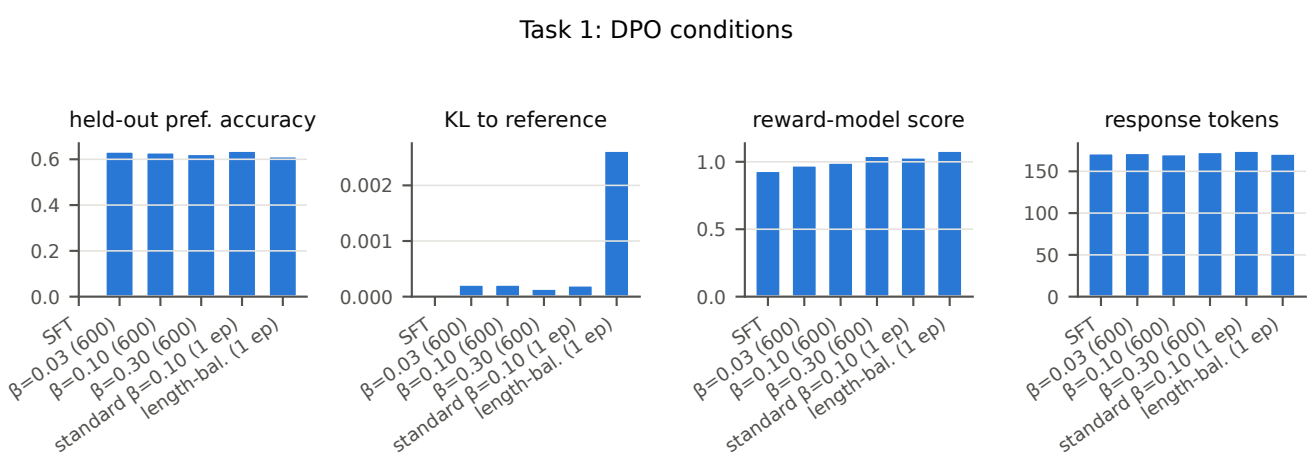
Task 1: DPO conditions
held-out pref. accuracy
KL to reference
reward-model score
response tokens
0.6
0.4
0.2
0.0
0.002
0.001
0.000
1.0
0.5
0.0
150
100
50
0
SFT
β=0.03 (600)
β=0.10 (600)
β=0.30 (600)
standard β=0.10 (1 ep)
length-bal. (1 ep)