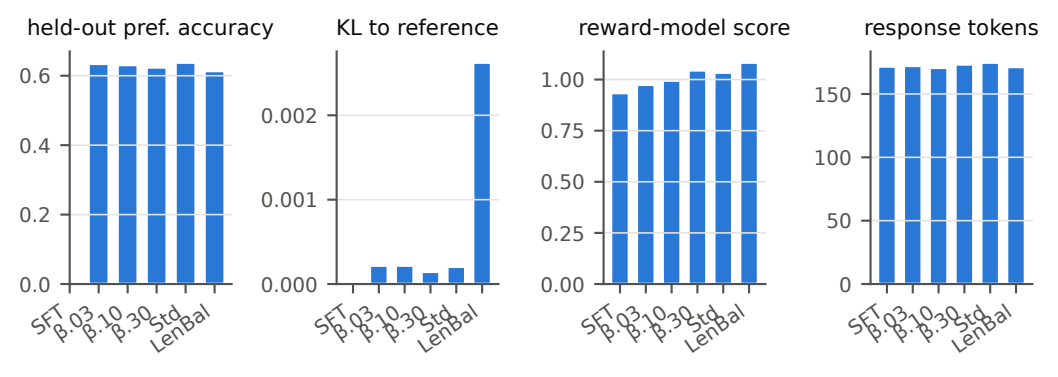
held-out pref. accuracy
0.6
0.4
0.2
0.0
KL to reference
0.002
0.001
0.000
reward-model score
1.00
0.75
0.50
0.25
0.00
response tokens
150
100
50
0
SFT
β.03
β.10
β.30
Std
LenBal