



Task 1: DPO conditions


held-out pref. accuracy




KL to reference



reward-model score



response tokens



SFT
 $\beta=0.03$ (600)
 $\beta=0.10$ (600)
 $\beta=0.30$ (600)
standard $\beta=0.10$ (1 ep)
length-bal. (1 ep)

SFT
 $\beta=0.03$ (600)
 $\beta=0.10$ (600)
 $\beta=0.30$ (600)
standard $\beta=0.10$ (1 ep)
length-bal. (1 ep)

SFT
 $\beta=0.03$ (600)
 $\beta=0.10$ (600)
 $\beta=0.30$ (600)
standard $\beta=0.10$ (1 ep)
length-bal. (1 ep)

SFT
 $\beta=0.03$ (600)
 $\beta=0.10$ (600)
 $\beta=0.30$ (600)
standard $\beta=0.10$ (1 ep)
length-bal. (1 ep)