



Task 2: standard PPO continuation

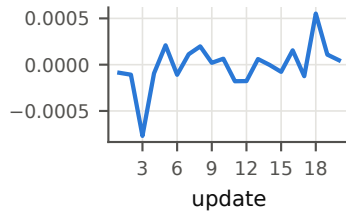
learned reward



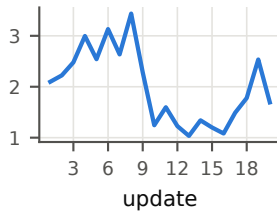
KL to reference



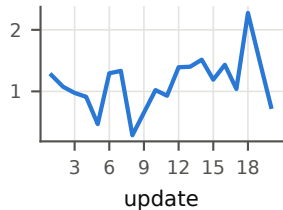
policy loss



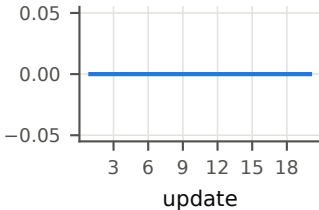
value loss



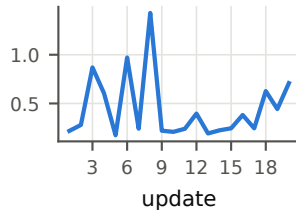
entropy



clip fraction



policy grad norm



response tokens

