



Task 3: standard GRPO continuation

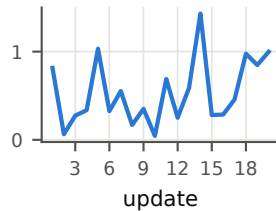
reward




KL to reference




within-group reward std




uninformative groups




policy loss



grad norm



entropy



response tokens

