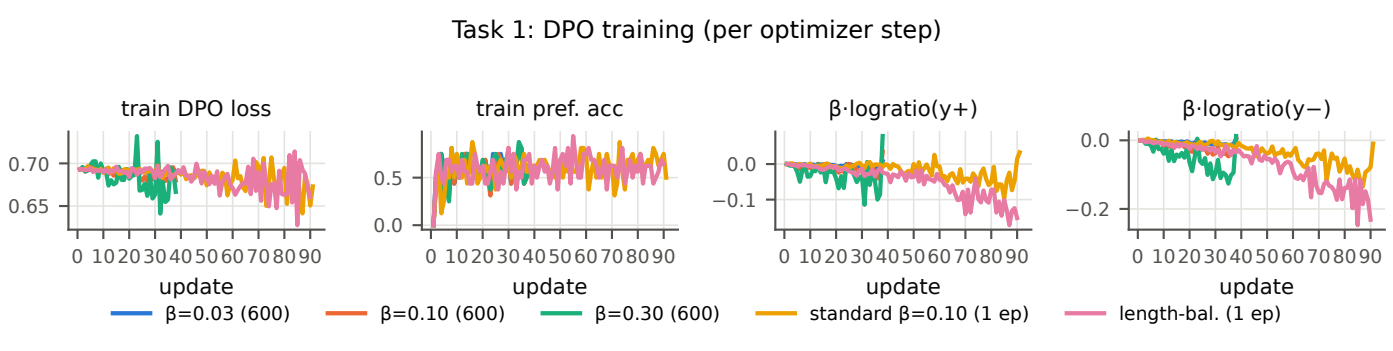
Task 1: DPO training (per optimizer step)
train DPO loss
0.70
0.65
0 10 20 30 40 50 60 70 80 90
update
train pref. acc
0.5
0.0
0 10 20 30 40 50 60 70 80 90
update
β·logratio(y+)
0.0
−0.1
0 10 20 30 40 50 60 70 80 90
update
β·logratio(y−)
0.0
−0.2
0 10 20 30 40 50 60 70 80 90
update
β=0.03 (600)
β=0.10 (600)
β=0.30 (600)
standard β=0.10 (1 ep)
length-bal. (1 ep)